

Introduction to Computer Networks



Reliable Transmission

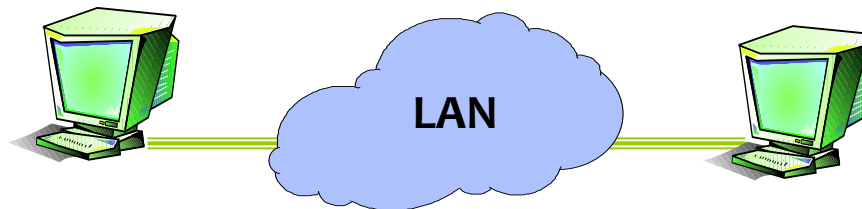
© All rights reserved. No part of this publication and file may be reproduced, stored in a retrieval system, or transmitted in any form or by any means, electronic, mechanical, photocopying, recording or otherwise, without prior written permission of Professor Nen-Fu Huang (E-mail: nfhuang@cs.nthu.edu.tw).

Outline

- **Introduction**
- **Stop-and-Wait protocol**
- **Sliding Window Protocol**
- **Issues with Sliding Window Protocol**

Reliable Transmission

- Data transmissions **over a communication link** may have errors caused by signal interference.
- Usually, CRC is used to detect errors.
- Some error codes are strong enough to correct errors but the overhead is typically too high.
- Corrupt frames must be discarded.
- For reliable transmission, we must recover from these discarded frames.



Reliable Transmission

- We can do this by using a combination of two fundamental mechanisms
 - **Acknowledgements** receiver收到，于是给一个ack的回应
 - **Timeouts** receiver根本没有收到，所以需要发送方有timeout机制
- An **acknowledgement** (**ACK** for short) is a small control frame that a protocol sends back to its peer saying that it has received the earlier frame.
 - A **control frame** is a frame with header only (no data).
- The receipt of an *acknowledgement* indicates to the sender that its frame was successfully delivered.

Reliable Transmission

- If the sender does not receive an *acknowledgment* after a reasonable amount of time, then it retransmits the original frame.
- The action of waiting a reasonable amount of time is called a *timeout*.
- The general strategy of using *acknowledgements* and *timeouts* to implement reliable delivery is sometimes called Automatic Repeat reQuest (ARQ).
timeout *ack*

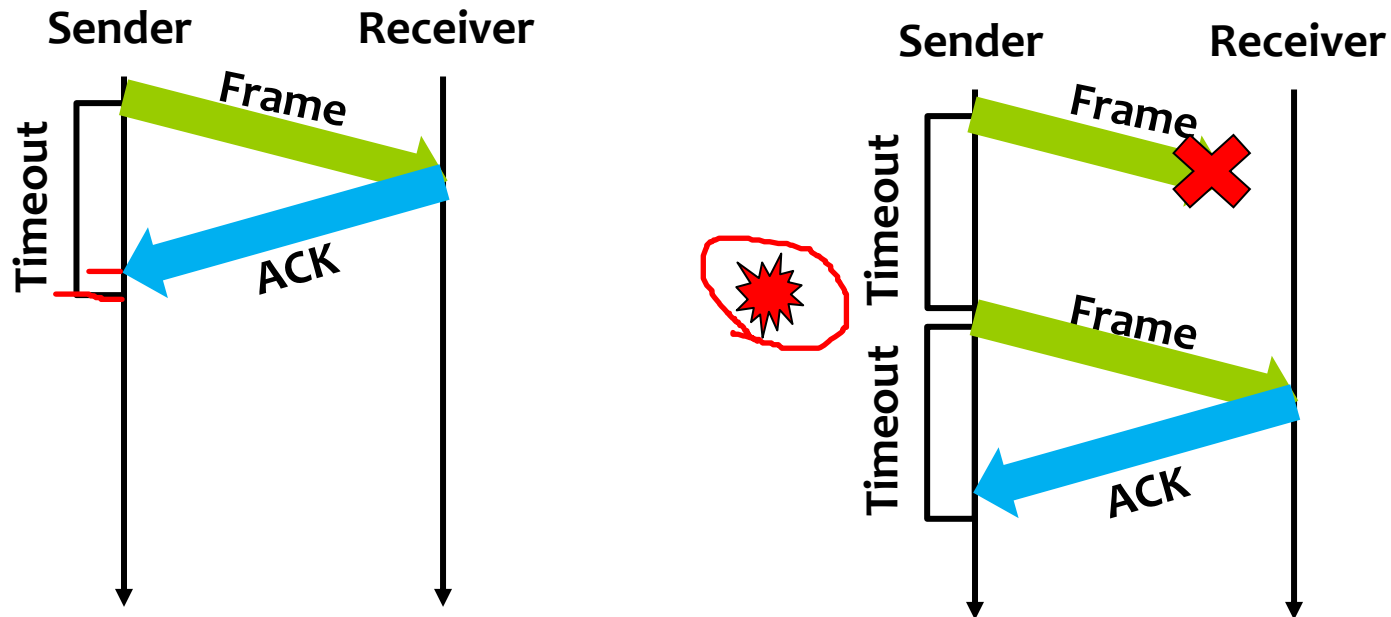
Outline

- Introduction
- **Stop-and-Wait protocol**
- Sliding Window Protocol
- Issues with Sliding Window Protocol

Stop and Wait Protocol

- Idea of **stop-and-wait protocol** is straightforward
送一个封包，等一个回应
- After transmitting one frame, the sender **waits for an acknowledgement** before transmitting the next frame.
- If the acknowledgement does not arrive after a certain period of time, the sender **times out and retransmits** the original frame

Stop and Wait Protocol

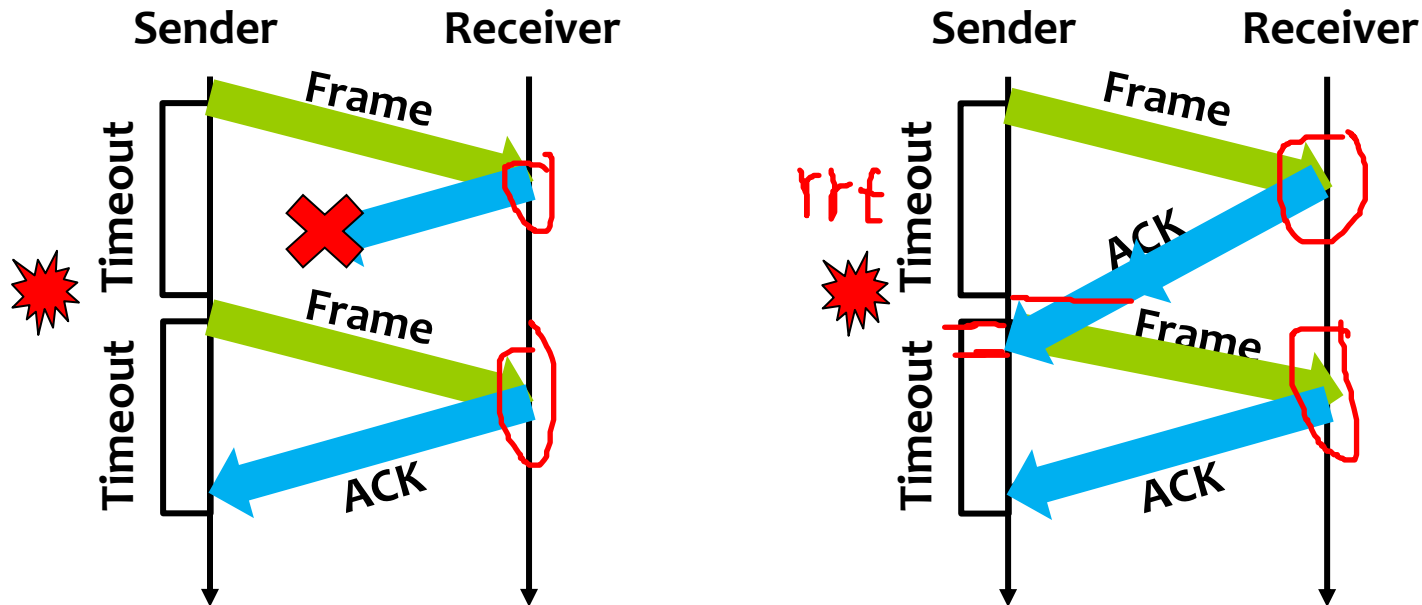


Four different scenarios for the stop-and-wait algorithm.

(a) The ACK is received before the timer expires;

(b) The original frame is lost;

Stop and Wait Protocol



Four different scenarios for the stop-and-wait algorithm.

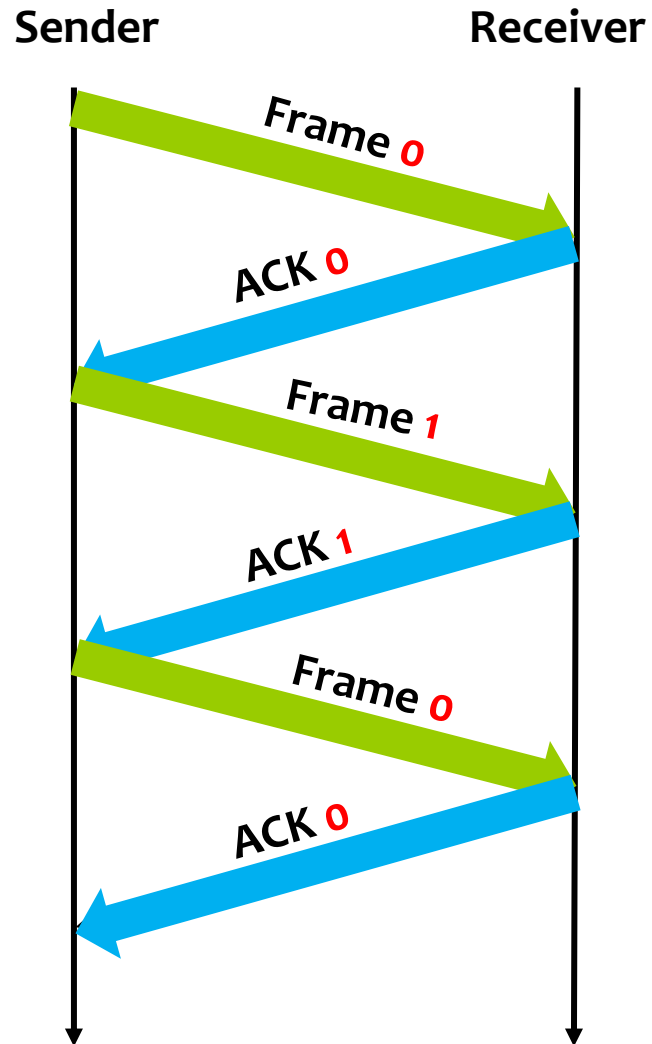
(c) The ACK is lost;

(d) The timeout fires too soon (or the ACK is delayed)

Stop and Wait Protocol

- If the acknowledgment is lost or delayed in arriving
 - The sender times out and retransmits the original frame.
 - As a result, duplicate copies of frames will be delivered
- How to solve this ?
 - Use 1 bit sequence number (0 or 1)
 - When the sender retransmits frame 0, the receiver can determine that this is a second copy of frame 0 rather than the first copy of frame 1 and therefore can ignore it (the receiver still acknowledges it, in case the first acknowledgement was lost)

Stop and Wait Protocol



Timeline for stop-and-wait with 1-bit sequence number


Stop and Wait Protocol

- The sender has **only one outstanding frame** on the link at a time -- Far below the link's capacity
- For example, consider a 2 Mbps link with a 40 ms RTT
 - The link has a delay \times bandwidth product of 80 Kb or **8 KB**
 - Since the sender can send only one frame per RTT and assuming **a frame size of 1 KB**
 - Maximum Sending rate
 - ▶ Bits per frame/Time per frame = $8\text{kb} \div 40\text{ms} = 200 \text{ Kbps}$
Or **about 1/10 of the link's capacity (2Mbps)**
 - To fully use the link, the sender should transmit up to ten frames before having to wait for an ACK.

Outline

- Introduction
- Stop-and-Wait protocol
- **Sliding Window Protocol**
- Issues with Sliding Window Protocol

Sliding Window Protocol



Timeline for Sliding Window Protocol


Sliding Window Protocol

- Sender assigns a sequence number denoted as **SeqNum** to each frame.
 - Assume it can grow infinitely large
- Sender maintains three variables
 - **Sending Window Size (SWS)**
 - ▶ Upper bound on the number of outstanding (unacknowledged) frames that the sender can transmit
 - **Last Acknowledgement Received (LAR)**
 - ▶ Sequence number of the last acknowledgement received
 - **Last Frame Sent (LFS)**
 - ▶ Sequence number of the last frame sent

Sliding Window Protocol

- Sender also maintains the following invariant

$$LFS - LAR \leq SWS$$




Sliding Window on Sender

Sliding Window Protocol

- When an acknowledgement arrives
 - the sender **moves LAR to right**, thereby allowing the sender to transmit another frame
- Also the sender associates a **timer** with each frame it transmits
 - It retransmits the frame if the timer expires before the ACK is received
- Note that the sender needs to **buffer** up to SWS frames for retransmissions, if necessary

Sliding Window Protocol

- When the sender is allowed to slide its Window ?
 - When the acknowledgement of $LAR+1$ is received



Sliding Window on Sender

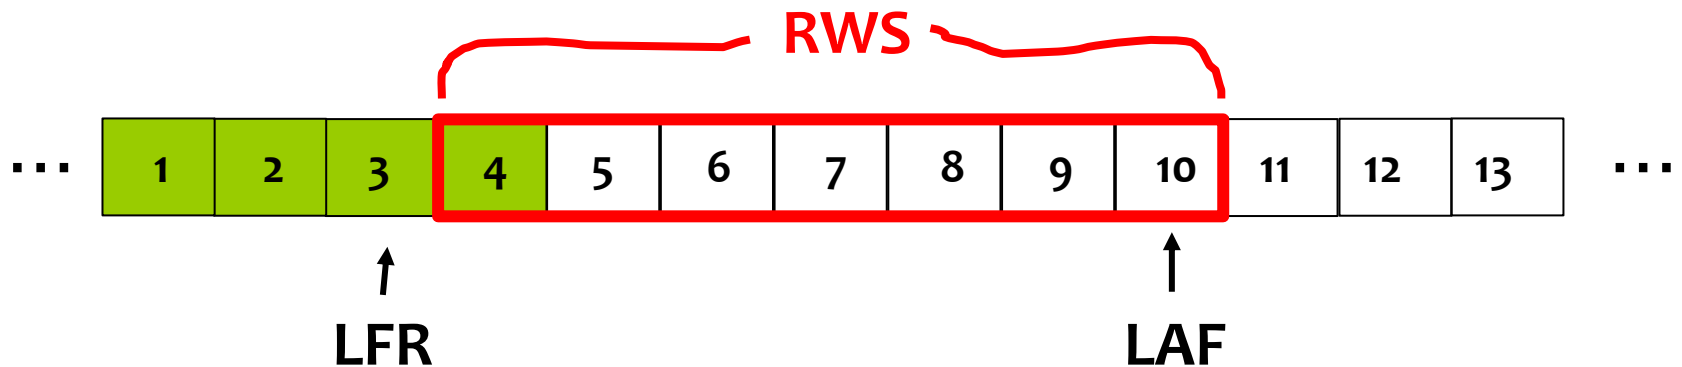
Sliding Window Protocol

- Receiver maintains three variables
 - Receiving Window Size (RWS)
 - ▶ Upper bound on the number of out-of-order frames that the receiver is willing to accept
 - Largest Acceptable Frame (LAF)
 - ▶ Sequence number of the largest acceptable frame
 - Last Frame Received (LFR)
 - ▶ Sequence number of the last frame received

Sliding Window Protocol

- Receiver also maintains the following invariant

$$\text{LAF} - \text{LFR} \leq \text{RWS}$$



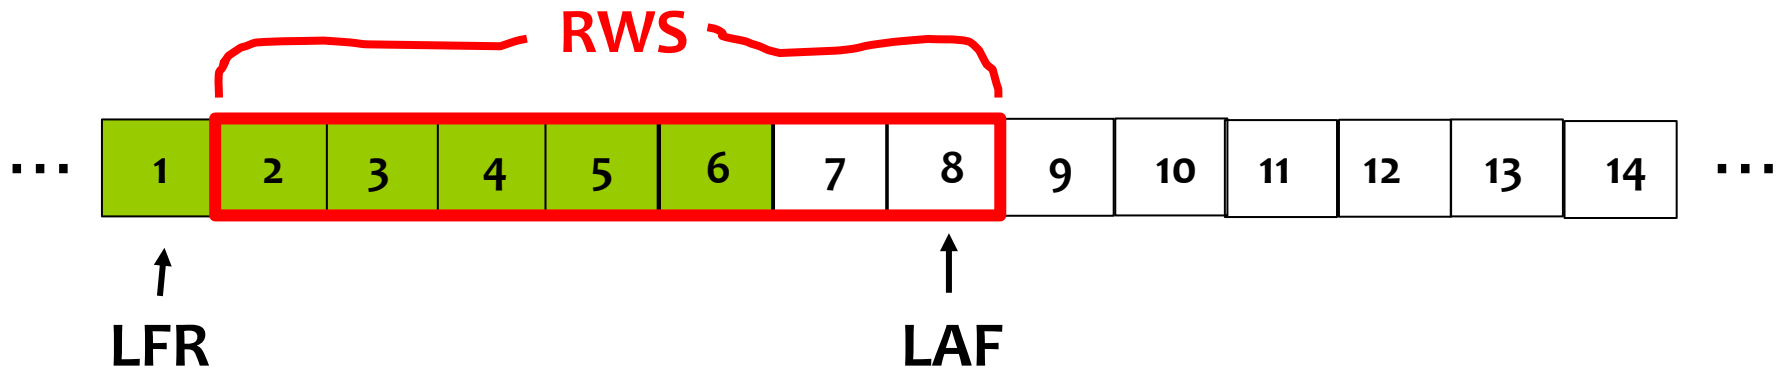
Sliding Window on Receiver

Sliding Window Protocol

- When a frame with sequence number **SeqNum** arrives, what does the receiver do?
 - If **$\text{SeqNum} \leq \text{LFR}$ or $\text{SeqNum} > \text{LAF}$**
 - ▶ Discard it (**the frame is outside the receiver window**)
 - If **$\text{LFR} < \text{SeqNum} \leq \text{LAF}$**
 - ▶ Accept it

Sliding Window Protocol

- When the receiver is allowed to slide its Window ?
 - When the frame of LFR+1 is received



Sliding Window on Receiver

Sliding Window Protocol


- For example, suppose **LFR = 1** and **RWS = 7**
 → LAF = 8
- If frames 4,6,3 and 5 arrive sequentially, they are out of order, but will still be buffered because they are within the receiver window
- But no ACK will be sent since frame 2 is yet to arrive
- Finally, frame 2 arrives (retransmitted or delayed)
- Now Receiver **Acknowledges Frame 6**
 and bumps LFR to 6
 and LAF to 13 (Window sliding)

Sliding Window Protocol

■ Cumulative Acknowledgement

- Let **SeqNumToAck** denote the largest sequence number not yet acknowledged, such that all frames with sequence number **less than SeqNumToAck have been received**
- The receiver acknowledges the receipt of **SeqNumToAck** even if high-numbered packets have been received
 - ▶ This acknowledgement is said to be **cumulative**.
- The receiver then sets
 - ▶ **$LFR = SeqNumToAck - 1$ and**
 - ▶ **$LAF = LFR + RWS$ (Window sliding)**

Sliding Window Protocol



Cumulative Acknowledgement example (**SeqNumToAck = 2**)

Outline

- Introduction
- Stop-and-Wait protocol
- Sliding Window Protocol
- Issues with Sliding Window Protocol

Issues with Sliding Window Protocol

- When **timeout occurs**, the amount of data in transit decreases
 - Since the sender is unable to advance its window
- When the **packet loss occurs**, this scheme is no longer keeping the pipe full
 - The longer it takes to notice that a packet loss has occurred, the more severe the problem becomes
- How to improve this
 - **Negative Acknowledgement (NAK)**
 - **Additional Acknowledgement**
 - **Selective Acknowledgement**

Issues with Sliding Window Protocol

■ Negative Acknowledgement (NAK)

- Receiver sends NAK for frame 2 when frame 3 arrive (in the cumulative acknowledgement example)
 - ▶ However this is unnecessary since sender's timeout mechanism will be sufficient to catch the situation

■ Additional Acknowledgement

- Receiver sends additional ACK for frame 2 when frame 3 arrives
 - ▶ Sender uses duplicate ACK as a clue for frame loss

■ Selective Acknowledgement

- Receiver will acknowledge exactly those frames it has received, rather than the highest number frames
 - ▶ Receiver will acknowledge frames 3,4,5, and 6
 - ▶ Sender knows frame 2 is lost
 - ▶ Sender can keep the pipe full (additional complexity)

Issues with Sliding Window Protocol

How to select the window size ?

- SWS is easy to calculate

- ▶ $\text{Delay} \times \text{Bandwidth}$

- RWS is more flexible

- ▶ Two common settings

- » $RWS = 1$

- No buffer at the receiver for frames that arrive out of order

- » $RWS = SWS$

- The receiver can buffer frames that the sender transmits

- Does it make any sense to keep $RWS > SWS$?

Issues with Sliding Window Protocol

■ Finite Sequence Number

- Frame sequence number is specified in the header field

- ▶ Finite size

- » 3 bits: eight possible sequence number: 0, 1, 2, 3, 4, 5, 6, 7

- ▶ It is necessary to wrap around

Issues with Sliding Window Protocol

- How to distinguish between different frames of the same sequence number?
 - Number of possible sequence number must be larger than the number of outstanding frames allowed
 - ▶ Stop and Wait: One outstanding frame
 - » 2 distinct sequence number (0 and 1)
 - ▶ Let **MaxSeqNum** be the number of available sequence numbers
 - ▶ $SWS + 1 \leq \text{MaxSeqNum}$
 - Is this sufficient?
 - Depends on RWS
 - If $RWS = 1$, then sufficient
 - If $RWS = SWS$, then not good enough

Issues with Sliding Window Protocol

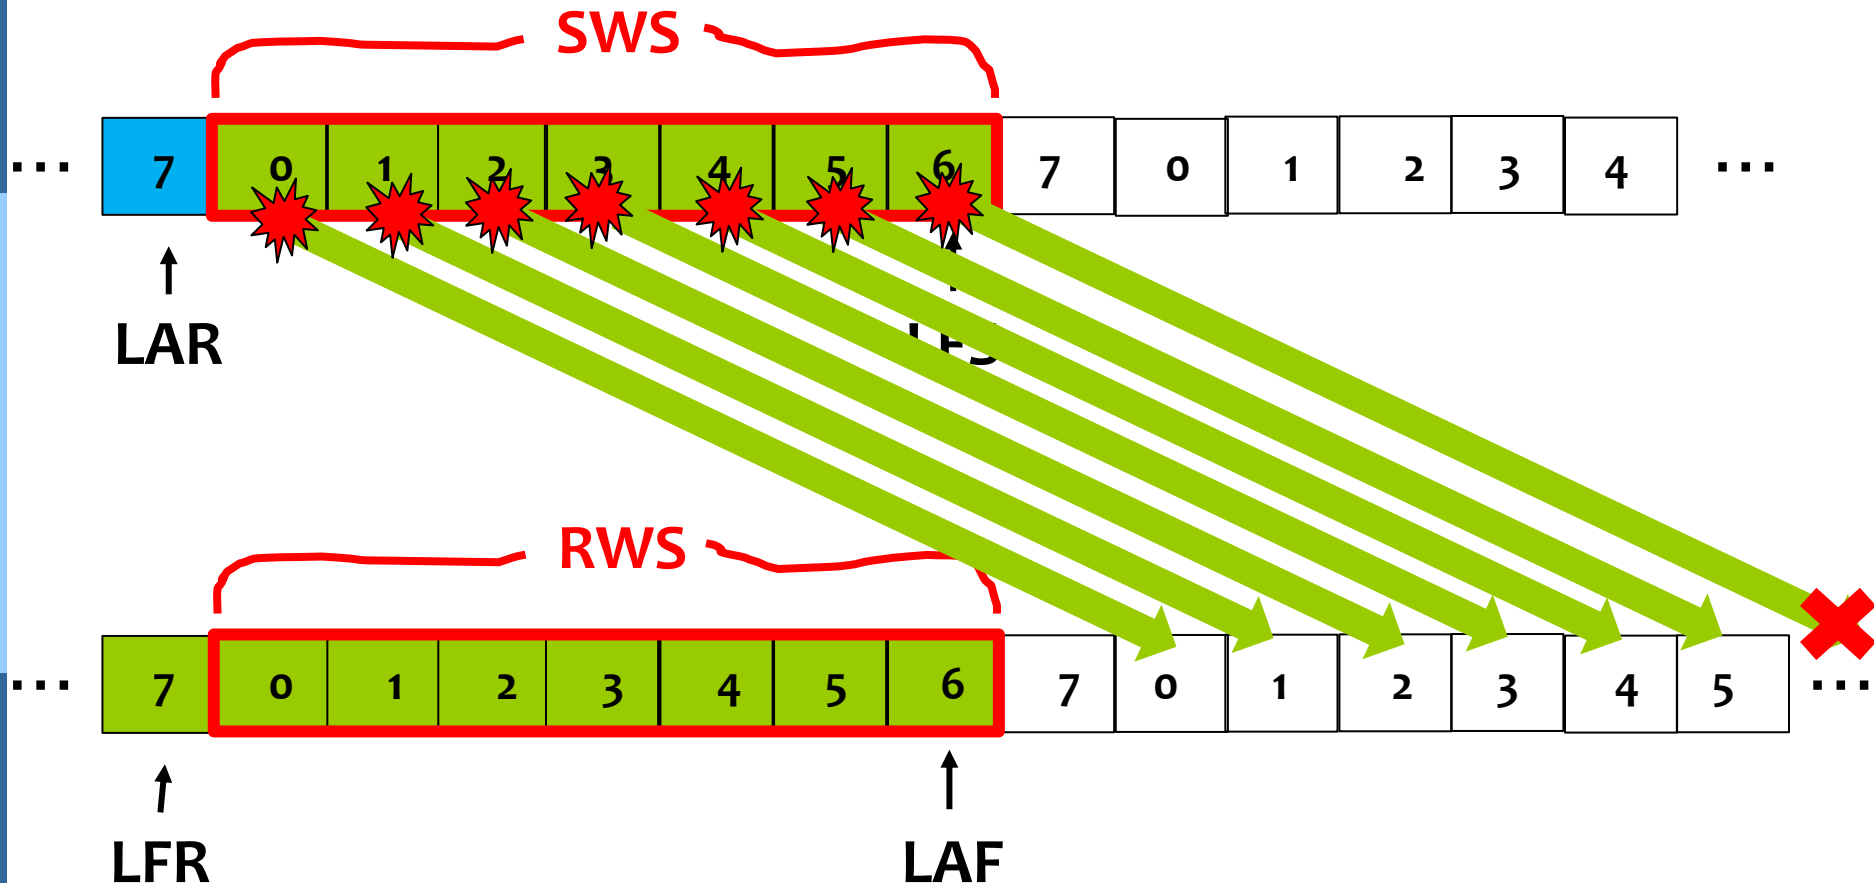
For example, we have eight sequence numbers

0, 1, 2, 3, 4, 5, 6, 7

$$RWS = SWS = 7 = 8-1$$

1. Sender sends 0, 1, ..., 6
2. Receiver receives 0, 1, ... ,6
3. Receiver acknowledges 0, 1, ..., 6
but **ACK (0, 1, ..., 6) are lost**
4. Sender timeouts and retransmits **0, 1, ...,5, 6**
5. Receiver is expecting 7, **0,, 5**
→ frames 0-5 will be accepted (duplicated !!)
frame 6 will be discarded (correct)

Issues with Sliding Window Protocol



Problem when $SWS + 1 \leq MaxSeqNum$ and $SWS = RWS$
 $MaxSeqNum = 8, SWS = RWS = 7$

Issues with Sliding Window Protocol

To avoid this,

If $RWS = SWS$

$$SWS < (MaxSeqNum + 1)/2$$

Issues with Sliding Window Protocol

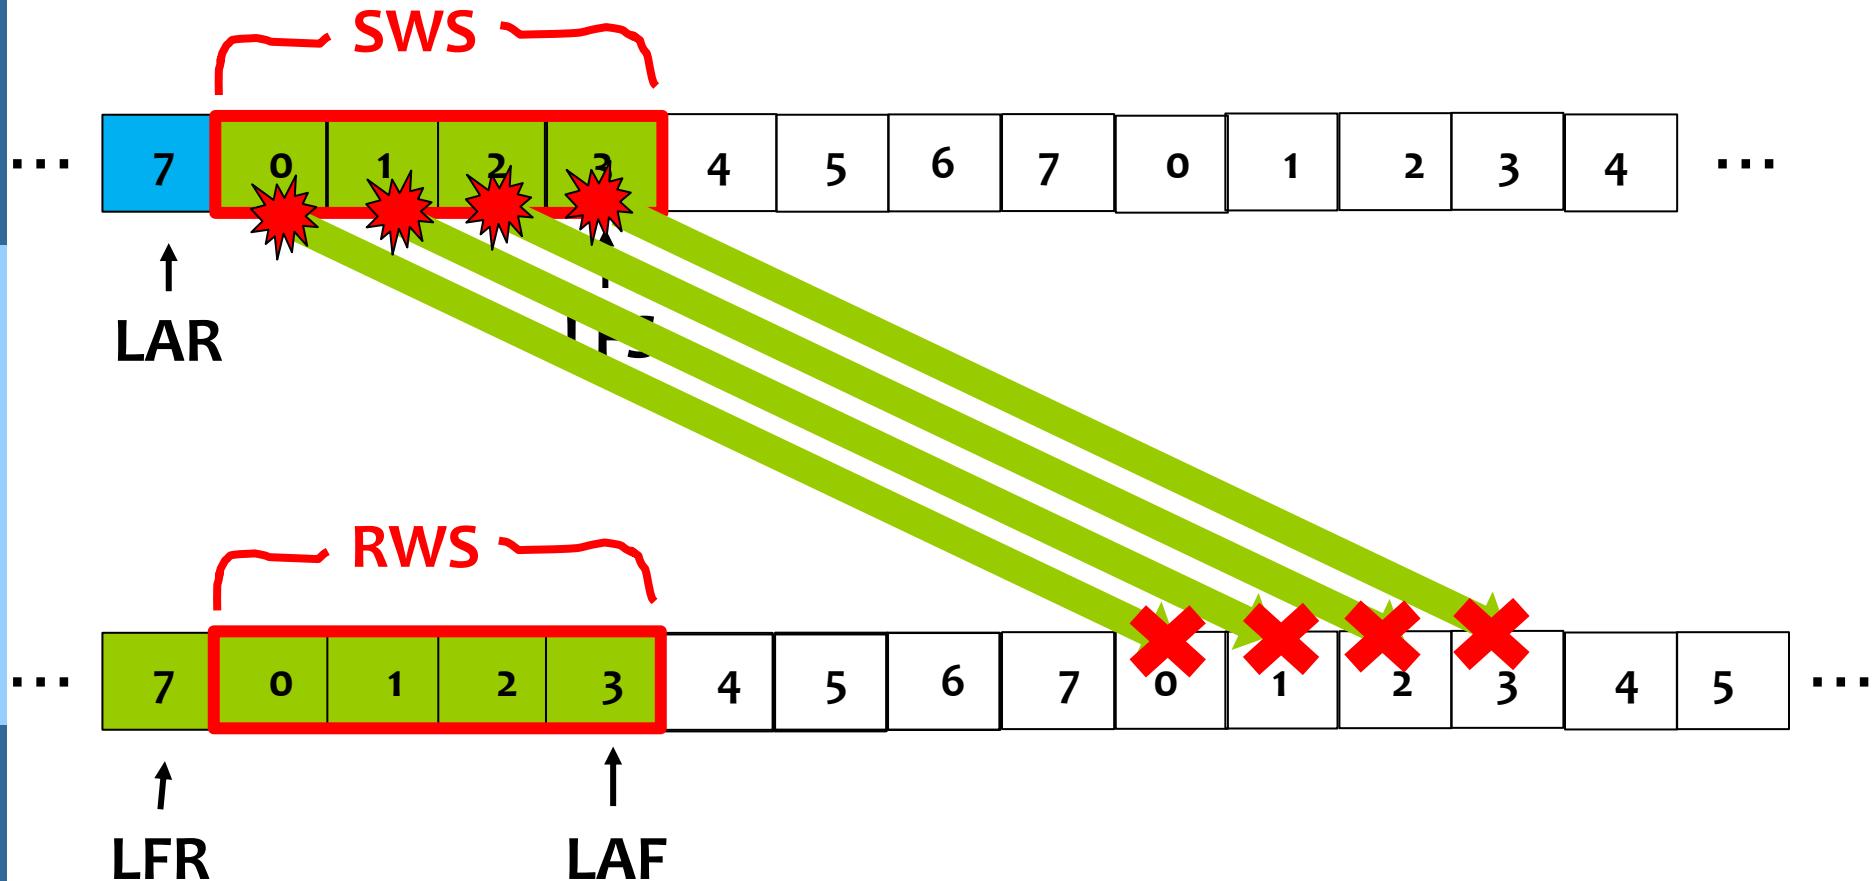
For example, we have eight sequence numbers

0, 1, 2, 3, 4, 5, 6, 7

$$RWS = SWS = 4 < (8+1)/2 = 4.5$$

1. Sender sends 0, 1, 2, 3
 2. Receiver receives 0, 1, 2, 3
 3. Receiver acknowledges 0, 1, 2, 3
but **ACK (0, 1, 2, 3) are lost**
 4. Sender timeouts and retransmits **0, 1, 2, 3**
 5. Receiver is expecting 4,5,6,7
- ➔ **all frames 0-3 will be discarded (correct now !!)**

Issues with Sliding Window Protocol



Example when $SWS < (MaxSeqNum+1)/2$
 $MaxSeqNum = 8, SWS = RWS = 4 < (8+1)/2$

Issues with Sliding Window Protocol

- Sliding Window Protocol provides three features
 - **Reliable Transmission**
 - **Preserve the order**
 - ▶ Each frame has a sequence number
 - ▶ Out of order frames will be buffered
 - **Flow control**
 - ▶ Receiver is able to throttle the sender by setting the value of RWS
 - ▶ Keeps the sender from overrunning the receiver

Summary

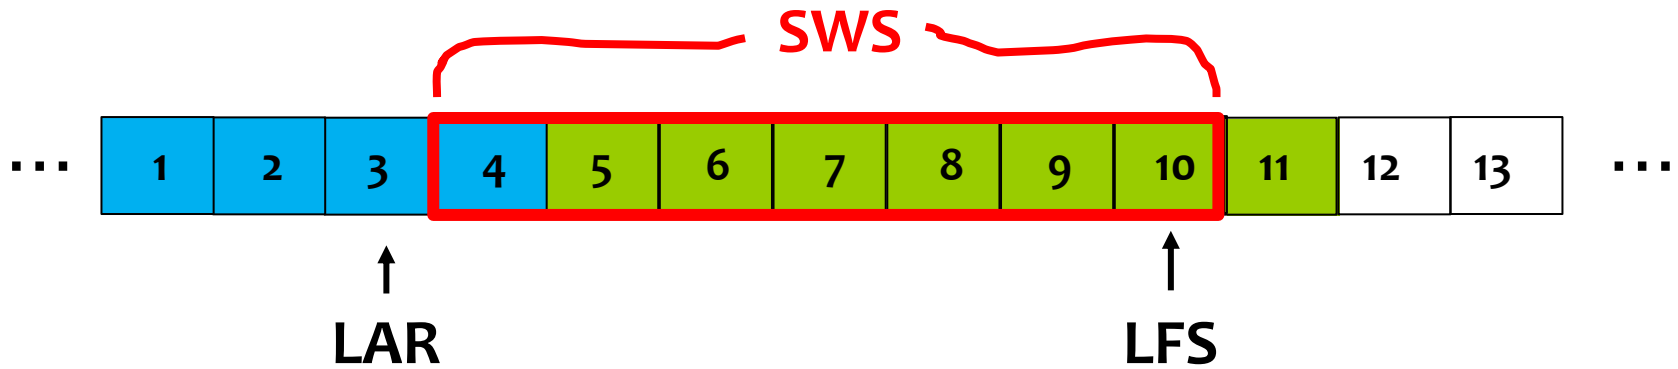
- Two fundamental mechanisms are used to provide reliable transmission over a communication link
 - Acknowledgements
 - Timeouts
- Stop-and-Wait Protocol is a reliable protocol, but the performance is not good enough
 - Only one outstanding frame
 - Receiver may receive duplicated frames

Summary

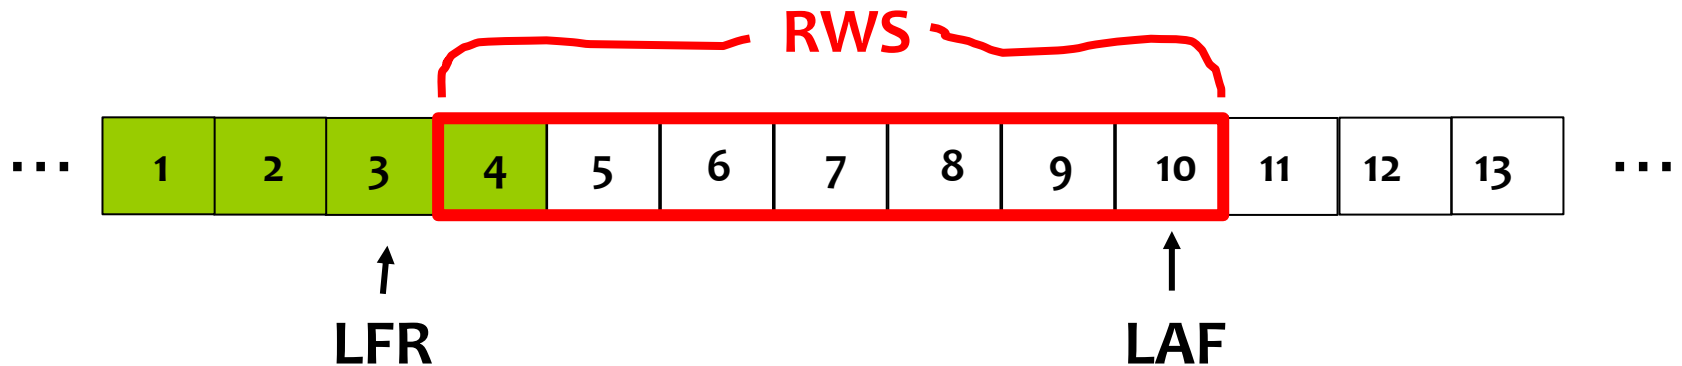
- **Sliding Window Protocol is a reliable and good performance protocol**
 - Sequence number is added for each frame
 - Multiple outstanding frames (keeping pipe full)
 - **Sender**
 - ▶ **Sending Window Size (SWS) = delay x bandwidth**
 - ▶ **Last Acknowledgement Received (LAR)**
 - ▶ **Last Frame Sent (LFS)**
 - **Receiver**
 - ▶ **Receiving Window Size (RWS) = 1 or SWS**
 - ▶ **Largest Acceptable Frame (LAF)**
 - ▶ **Last Frame Received (LFR)**

Summary

- Sender maintains the invariant: $LFS - LAR \leq SWS$



- Receiver maintains the invariant: $LAF - LFR \leq RWS$



Summary

- The values of SWS and RWS may be different.
- If $RWS = SWS$, it is better $SWS < (MaxSeqNum + 1)/2$ otherwise, the receiver may still receive duplicated frames.
- The Acknowledgement mechanism is an option
 - NAK
 - Cumulative acknowledgement
- Sliding Window Protocol provides three features
 - Reliable Transmission
 - Preserve the order
 - Flow control (receiver determines the RWS)